

---

**Dhirubhai Ambani University (DAU)**  
(Formerly Dhirubhai Ambani Institute of Information and Communication Technology)

Gandhinagar, Gujarat, India

---

# CS 301

## High-Performance Computing

---

### Lab 02: Matrix Multiplication

### Performance Analysis

**Submitted by:**

Ramani Drahosit (202301404)  
Vasu Bhesandariya (202301416)

**Department:**

Information and  
Communication Technology

February 17, 2026

## Contents

<b>1 Matrix Multiplication: Problem Size vs Time</b>	<b>2</b>
1.1 Standard Loop Permutations . . . . .	2
1.1.1 ijk Order . . . . .	2
1.1.2 ikj Order . . . . .	2
1.1.3 jik Order . . . . .	3
1.1.4 jki Order . . . . .	3
1.1.5 kij Order . . . . .	4
1.1.6 kji Order . . . . .	4
1.2 Optimized Matrix Multiplication . . . . .	5
1.2.1 Block Matrix Multiplication . . . . .	5
1.2.2 Transpose-based Matrix Multiplication . . . . .	5
<b>2 Conclusion</b>	<b>6</b>

# 1 Matrix Multiplication: Problem Size vs Time

## 1.1 Standard Loop Permutations

### 1.1.1 ijk Order

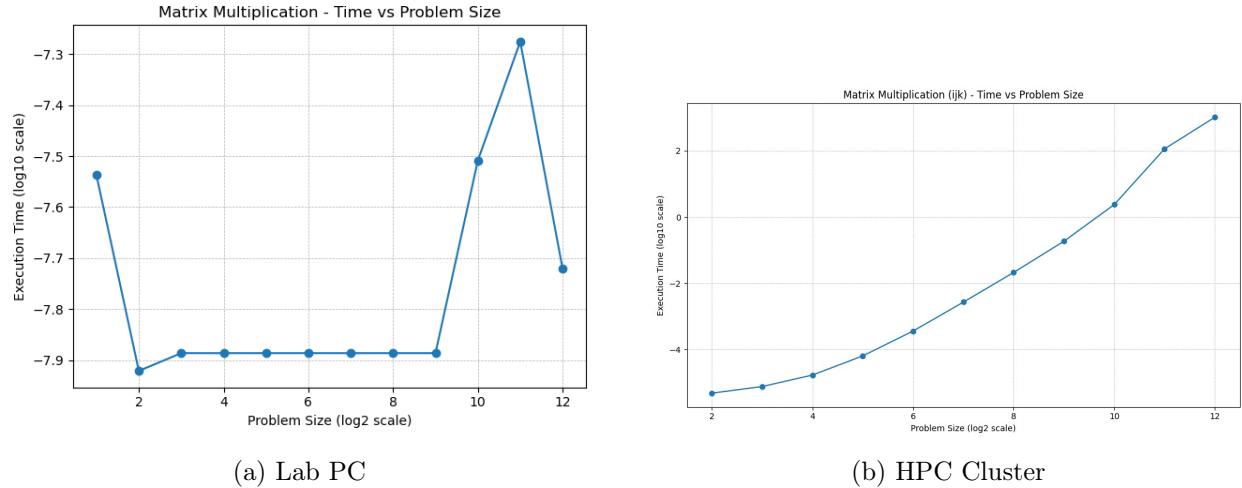


Figure 1: Problem Size vs Time (ijk)

### 1.1.2 ikj Order

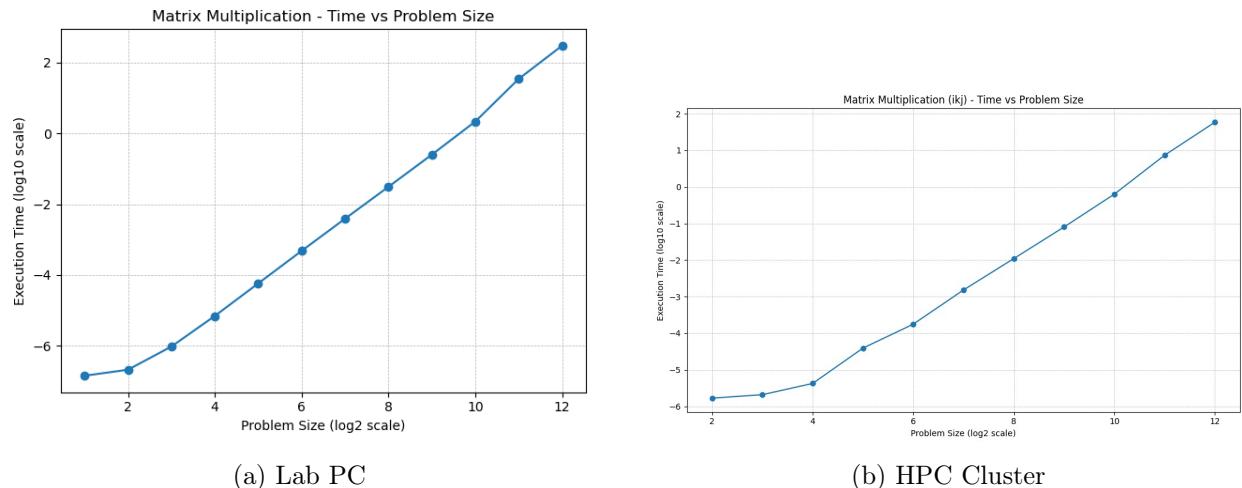


Figure 2: Problem Size vs Time (ikj)

### 1.1.3 jik Order

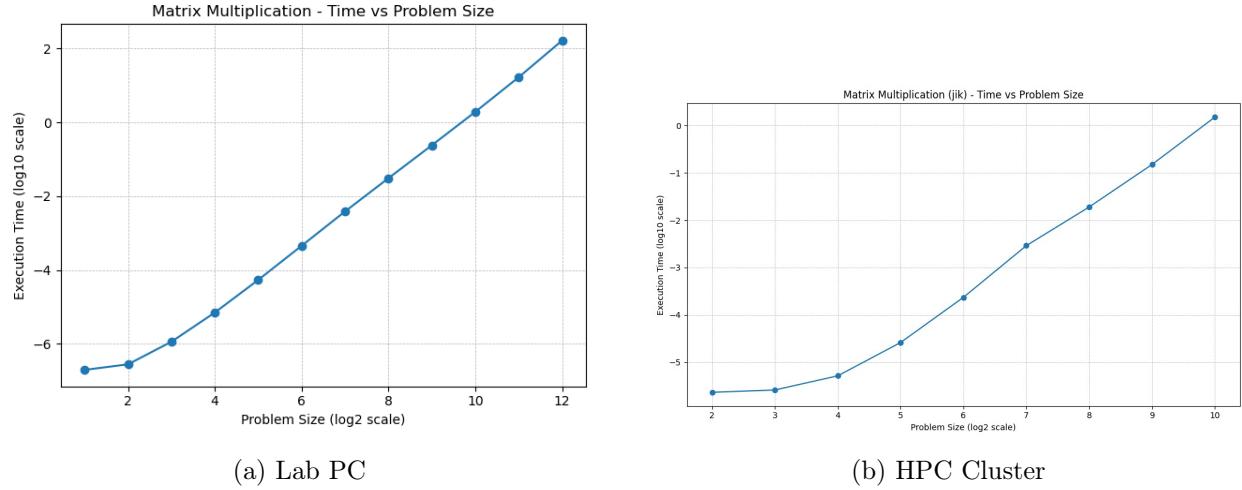


Figure 3: Problem Size vs Time (jik)

### 1.1.4 jki Order

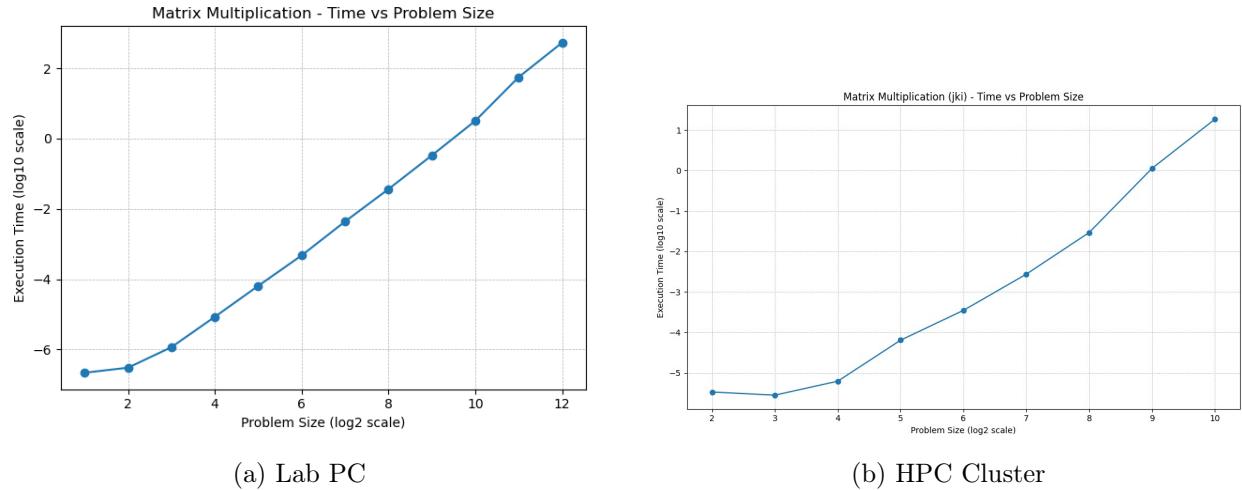


Figure 4: Problem Size vs Time (jki)

### 1.1.5 kij Order

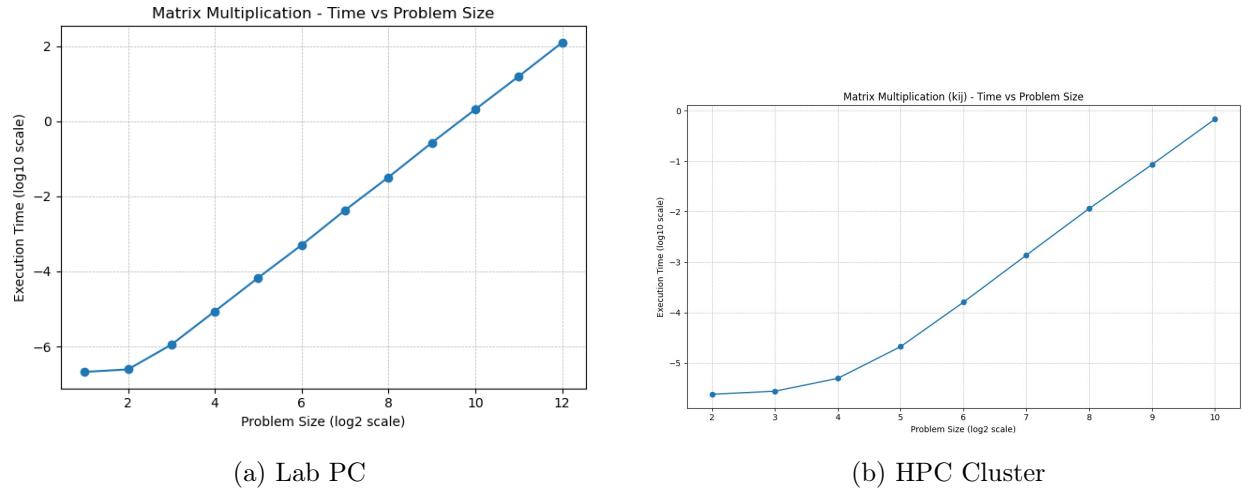


Figure 5: Problem Size vs Time (kij)

### 1.1.6 kji Order

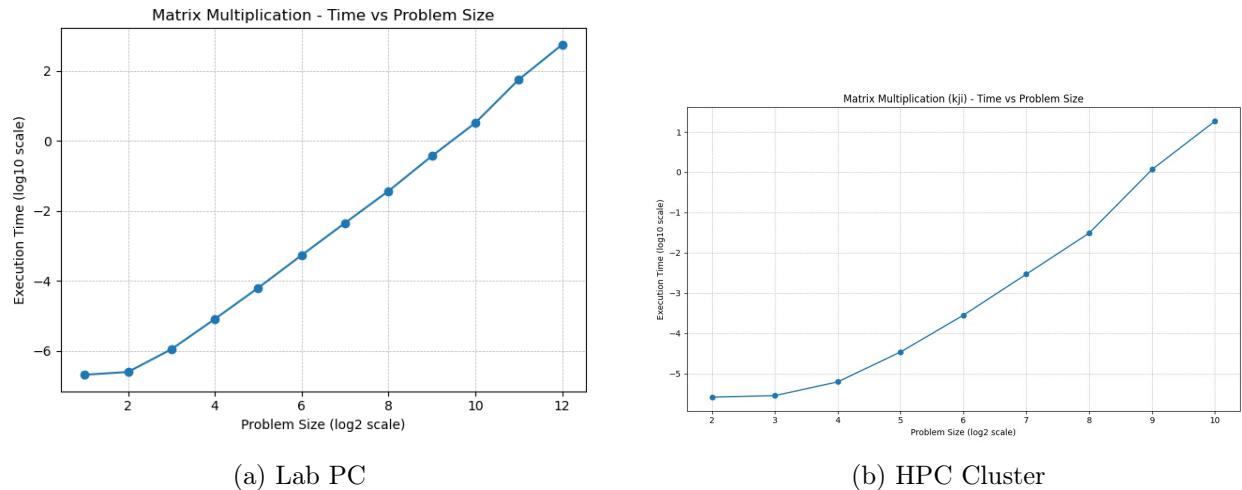


Figure 6: Problem Size vs Time (kji)

## 1.2 Optimized Matrix Multiplication

### 1.2.1 Block Matrix Multiplication

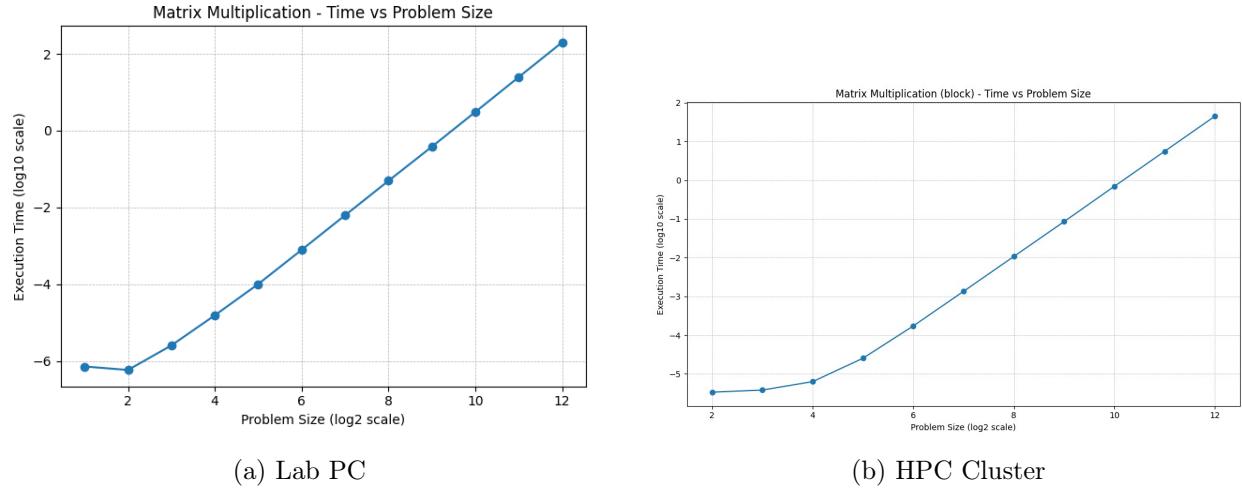


Figure 7: Problem Size vs Time (Block)

### 1.2.2 Transpose-based Matrix Multiplication

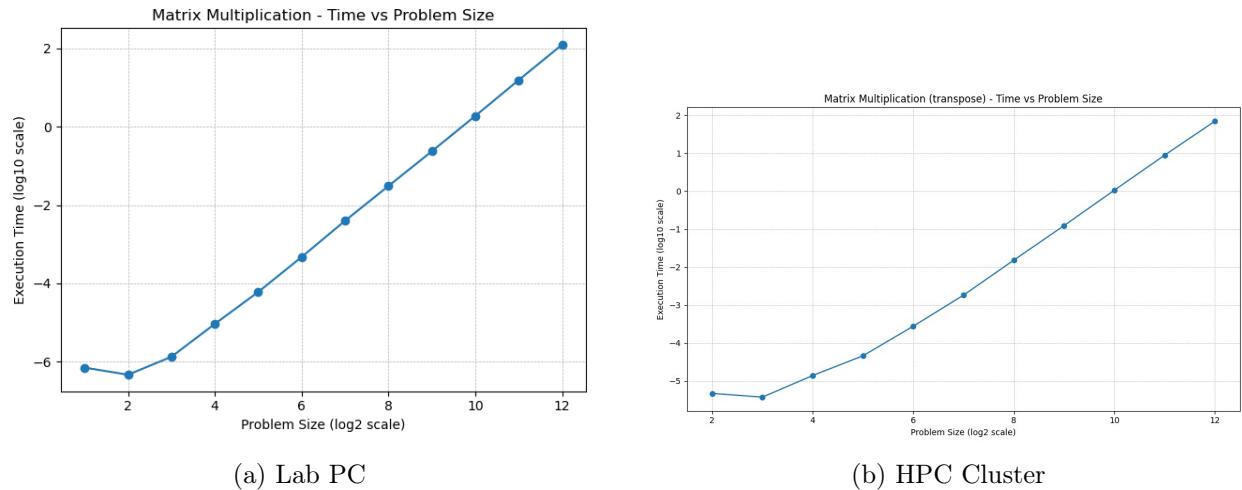


Figure 8: Problem Size vs Time (Transpose)

## 2 Conclusion

This report presented a performance analysis of various matrix multiplication algorithms. We observed significant differences in execution time across different loop orderings, as well as with the transpose and block optimization strategies.

Generally, algorithms that exhibit better cache locality (e.g., **ikj**, **kij**, **block**, and **transpose**) tend to perform significantly better than those with poor cache utilization (e.g., **ijk**, **jik**, **JKI**, and **kji**), especially as the problem size increases.

- **Memory Access Patterns:** The superior performance of the **block** matrix multiplication and the **transpose** method highlights the critical importance of spatial and temporal locality.
- **Hardware Impact:** While the Lab PC and HPC Cluster showed different absolute execution times, the trend remains consistent: hardware performance is heavily bottlenecked by memory bandwidth and cache misses rather than raw CPU cycles.
- **Optimization:** Effective cache usage is the primary driver for achieving high performance in numerical linear algebra.

The experimental results confirm that software-level optimizations targeting the memory hierarchy are essential for high-performance computing tasks.